

Project #2 PL/SQL (Procedures and Triggers)

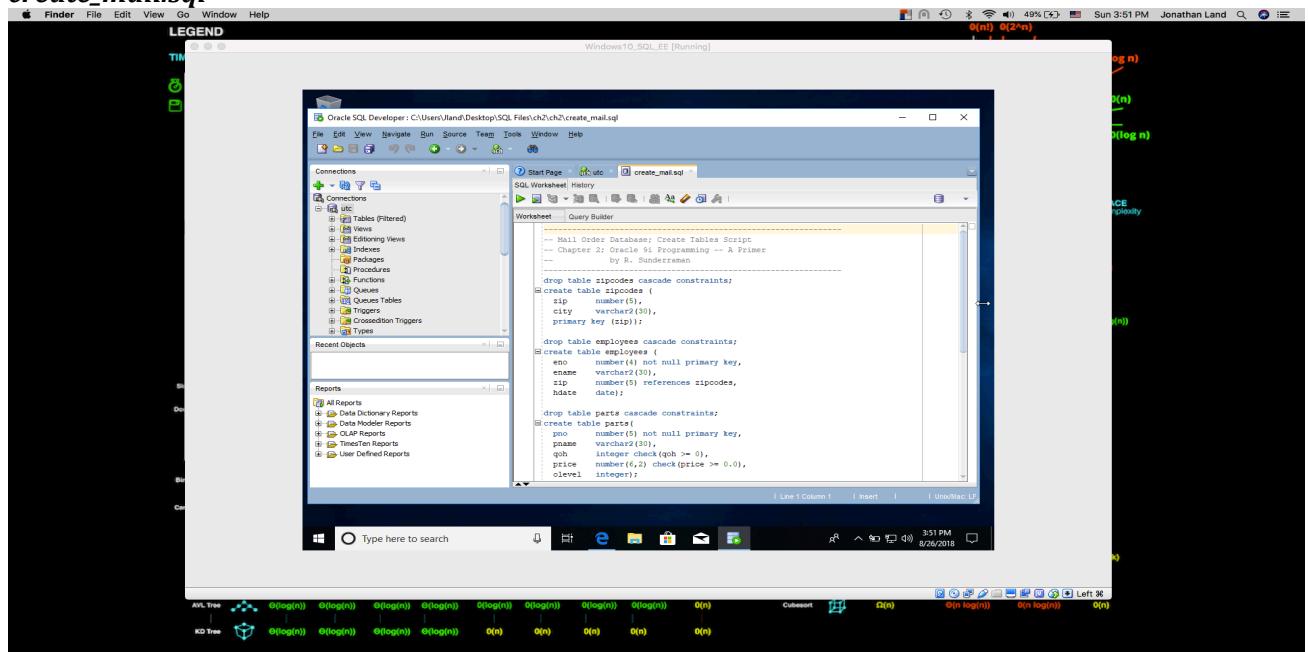
CPSC 5270 Advanced Database and Database Security

Jonathan Land

Preparation (20 points):

1. Execute scripts such as *create_mail.sql*, *insert_mail.sql*, *create_grades.sql*, and *insert_grades.sql* scripts in the zip file of *plssql.ch2.zip*.

create_mail.sql



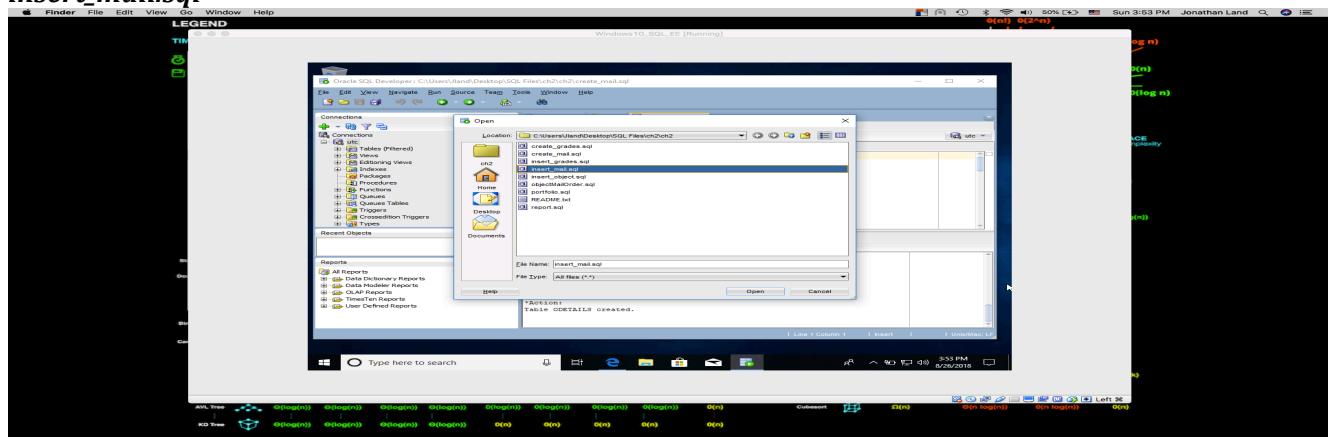
The screenshot shows the Oracle SQL Developer interface. The 'Worksheet' tab is active, displaying the contents of the 'create_mail.sql' script. The script creates three tables: 'zipcodes', 'employees', and 'parts'. It includes constraints like cascade delete for foreign keys and primary key definitions. The 'Connections' sidebar on the left shows a single connection named 'JLc'. The status bar at the bottom right indicates the session is running on 'Windows10_SQL_EE [Running]'.

```
-- Mail Order Database: Create Tables Script
-- Chapter 2: Oracle 9i Programming -- A Primer
-- by R. Sundaram

drop table zipcodes cascade constraints;
create table zipcodes (
    zip      number(5),
    city     varchar2(30),
    state   varchar2(2),
    primary key (zip));
    
drop table employees cascade constraints;
create table employees (
    empno  number(4) not null primary key,
    ename   varchar2(30),
    zip     number(5) references zipcodes,
    hiredate date);

drop table parts cascade constraints;
create table parts(
    pno    number(5) not null primary key,
    pname   varchar2(30),
    qoh    integer check(qoh >= 0),
    price   number(4,2) check(price >= 0.0),
    clevel  integer);
```

insert_mail.sql



insert_mail.sql

```

-- Mail Order Database: Insert Rows
-- Chapter 2: Oracle 9i Programming -- A Primer
-- by R. Sundaraman

insert into zipcodes values
(67226,'Wichita');
insert into zipcodes values
(60606,'Fort Dodge');
insert into zipcodes values
(50302,'Kansas City');

```

Script Output:

```

Task completed in 1.344 seconds

1 row inserted.

1 row inserted.

1 row inserted.

```

create_grades.sql

```

-- Grade Book Database: Create Tables Script
-- Chapter 2: Oracle 9i Programming -- A Primer
-- by R. Sundaraman

drop table catalog cascade constraints;
create table catalog (
  cno varchar2(7) not null,
  cttitle varchar2(50),
  primary key (cno);

```

Script Output:

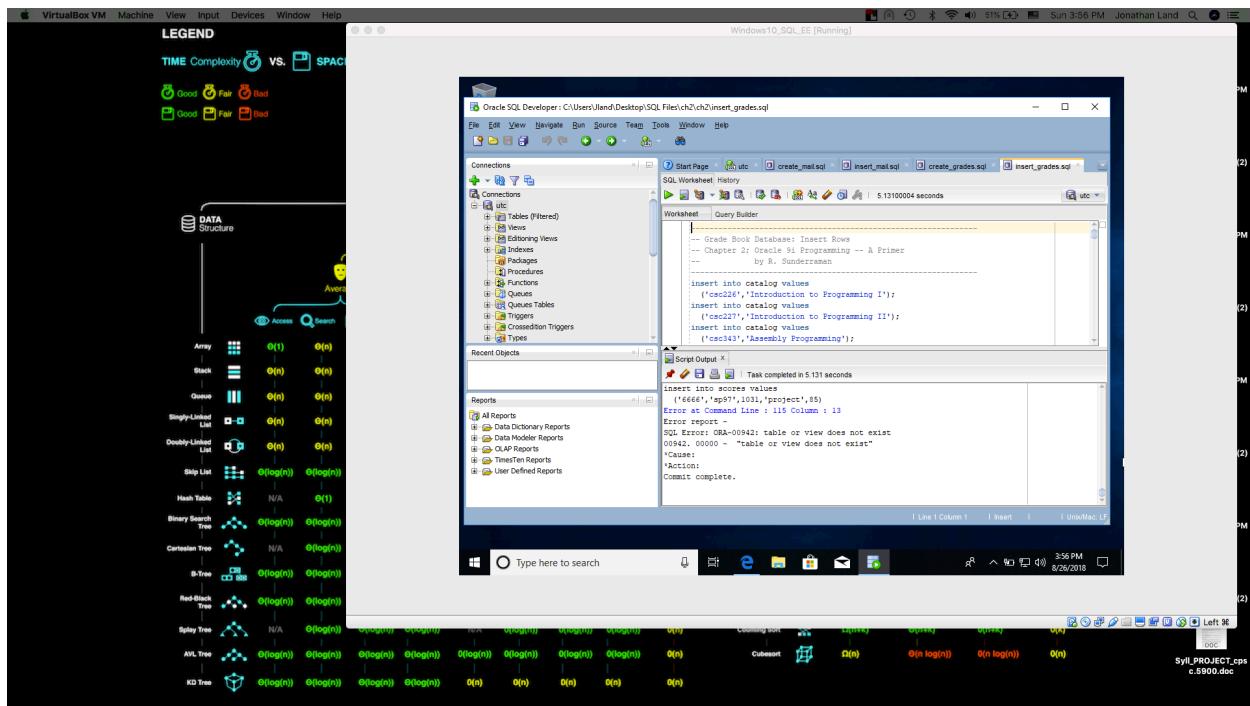
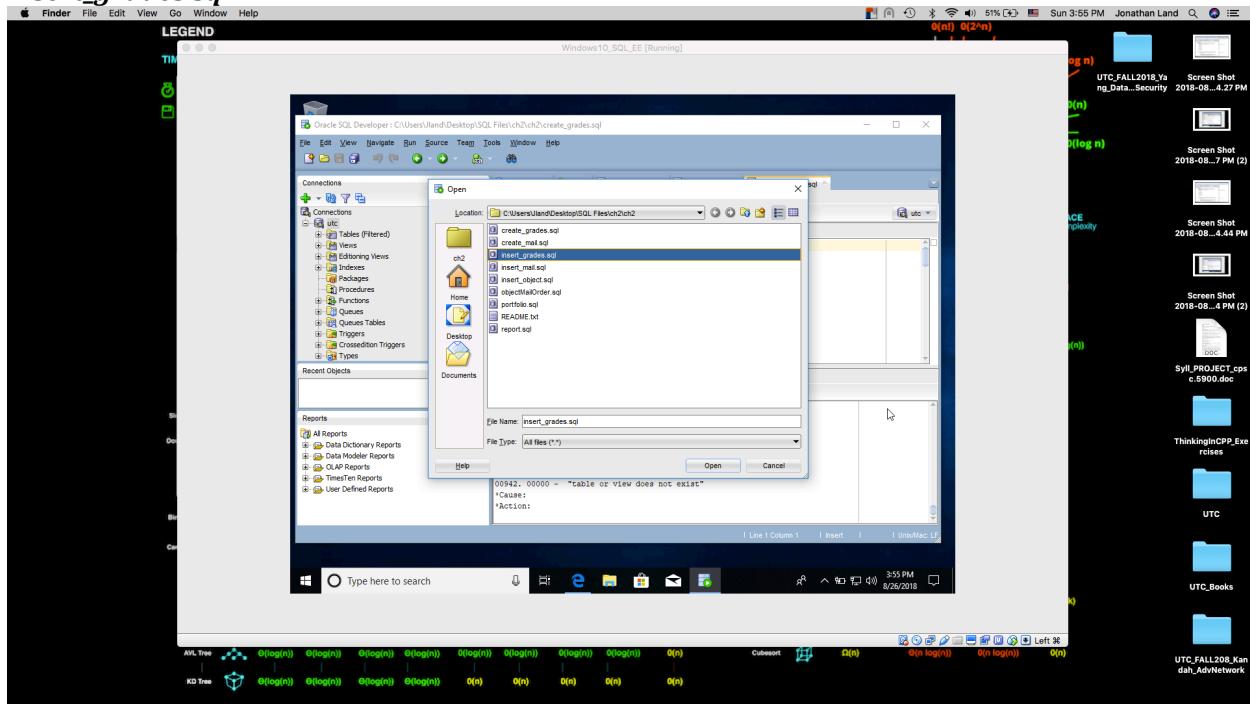
```

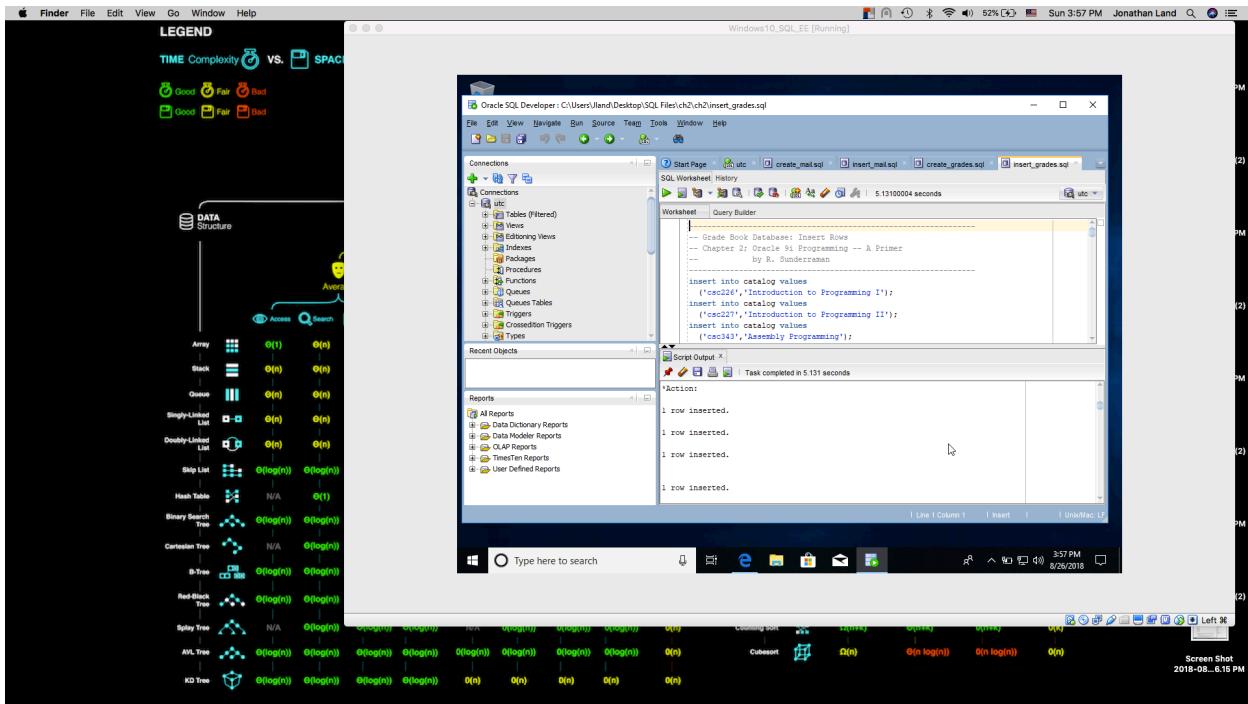
Task completed in 1.632 seconds

comname varchar2(115) not null,
points number(2,1) check(points >= 0),
primary key (sid,term,lineno,comname),
foreign key (sid,term,lineno) references enrolls,
foreign key (term,lineno,comname) references components)
Error report -
SQL Error: ORA-00942: table or view does not exist
00942. 00000 - "table or view does not exist"
*Cause:
*Action:

```

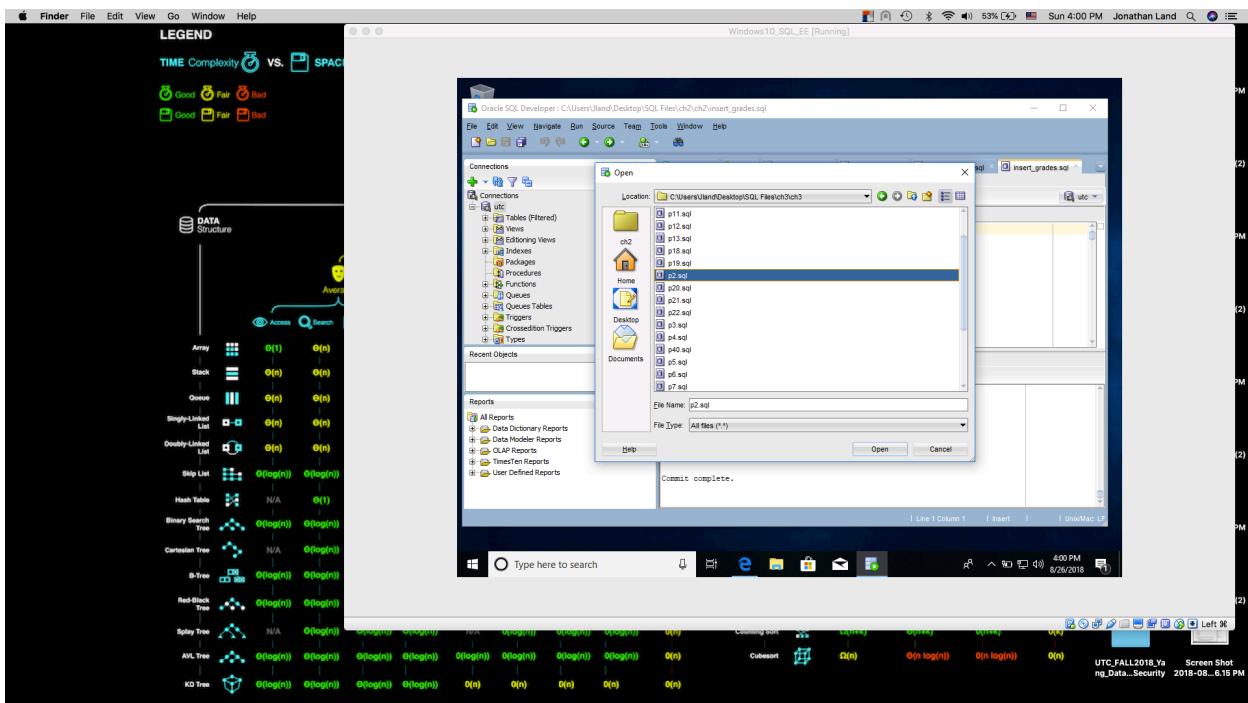
insert_grades.sql

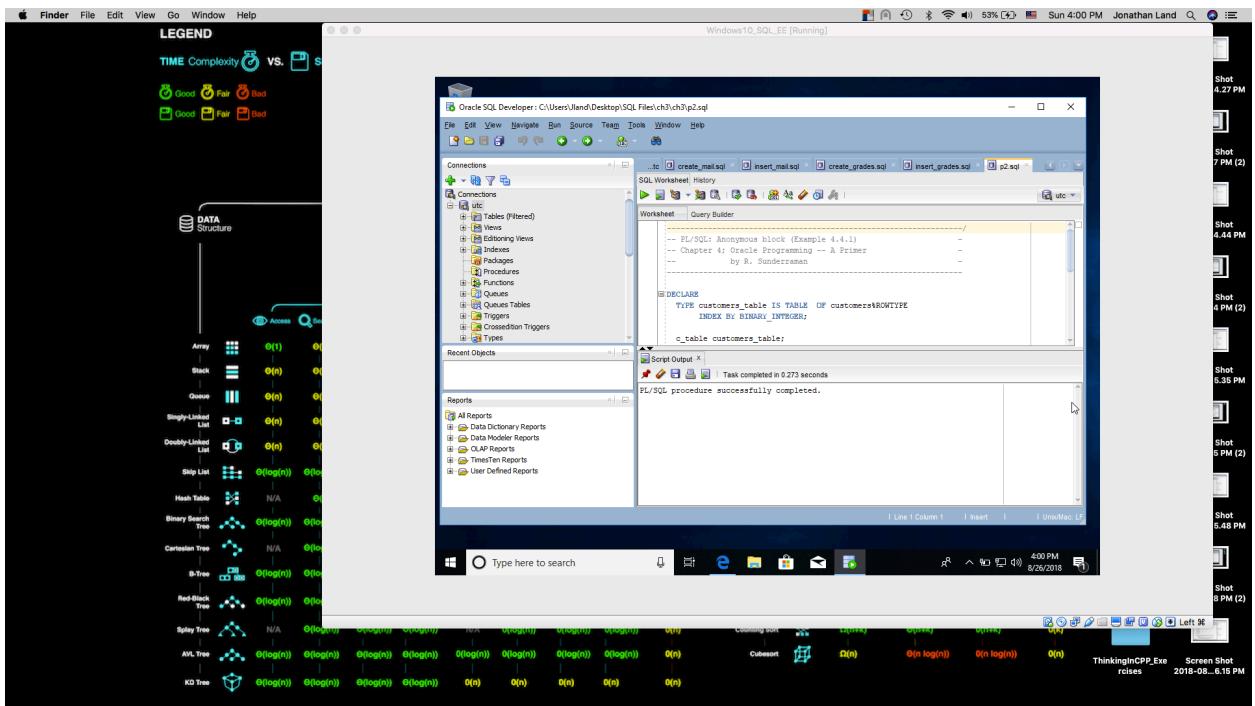
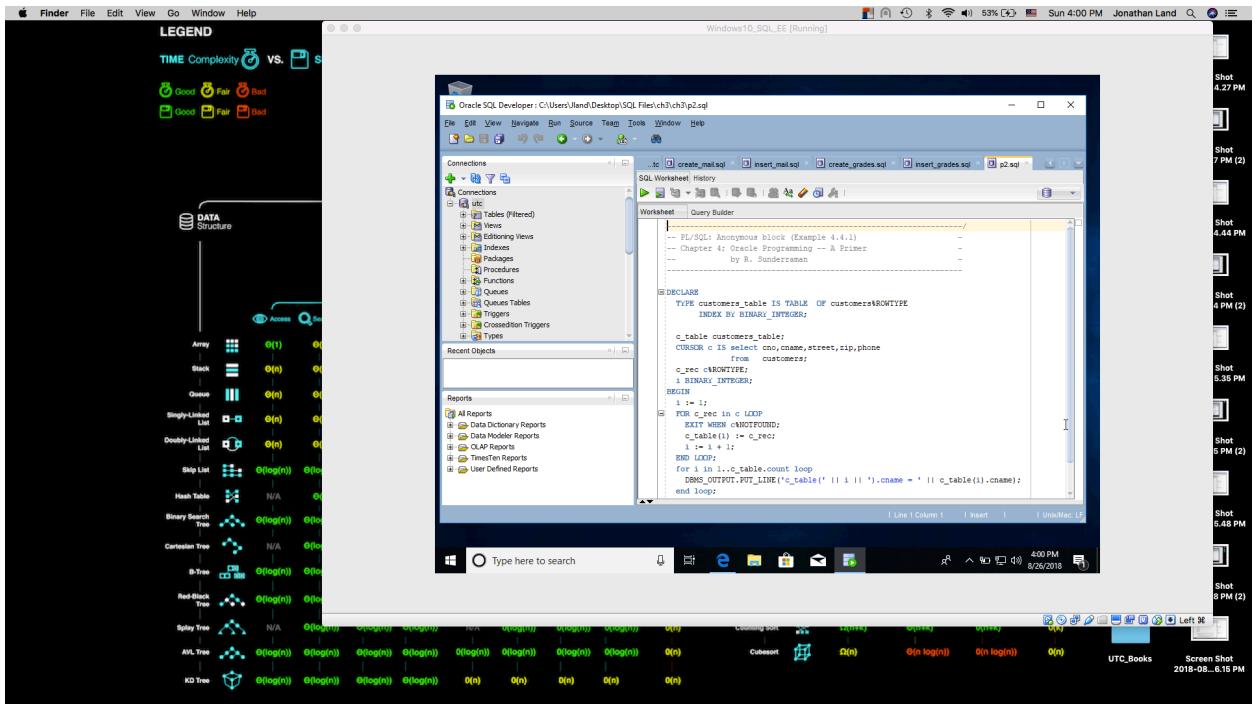




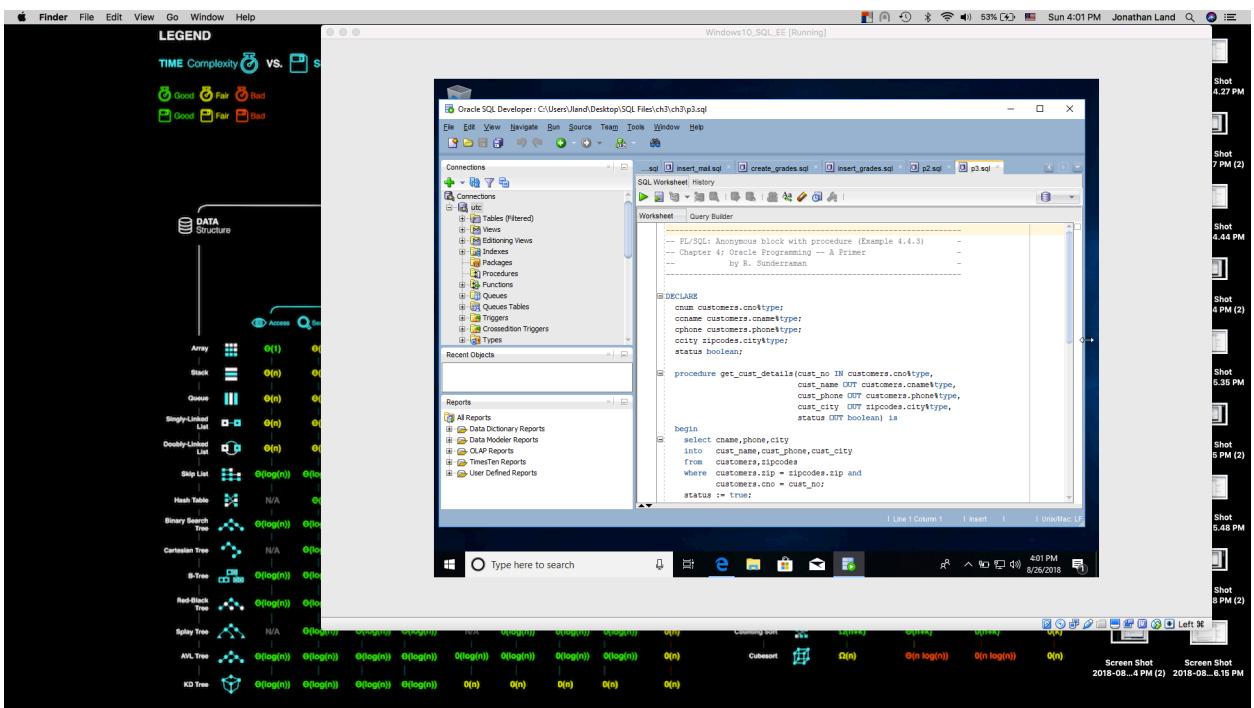
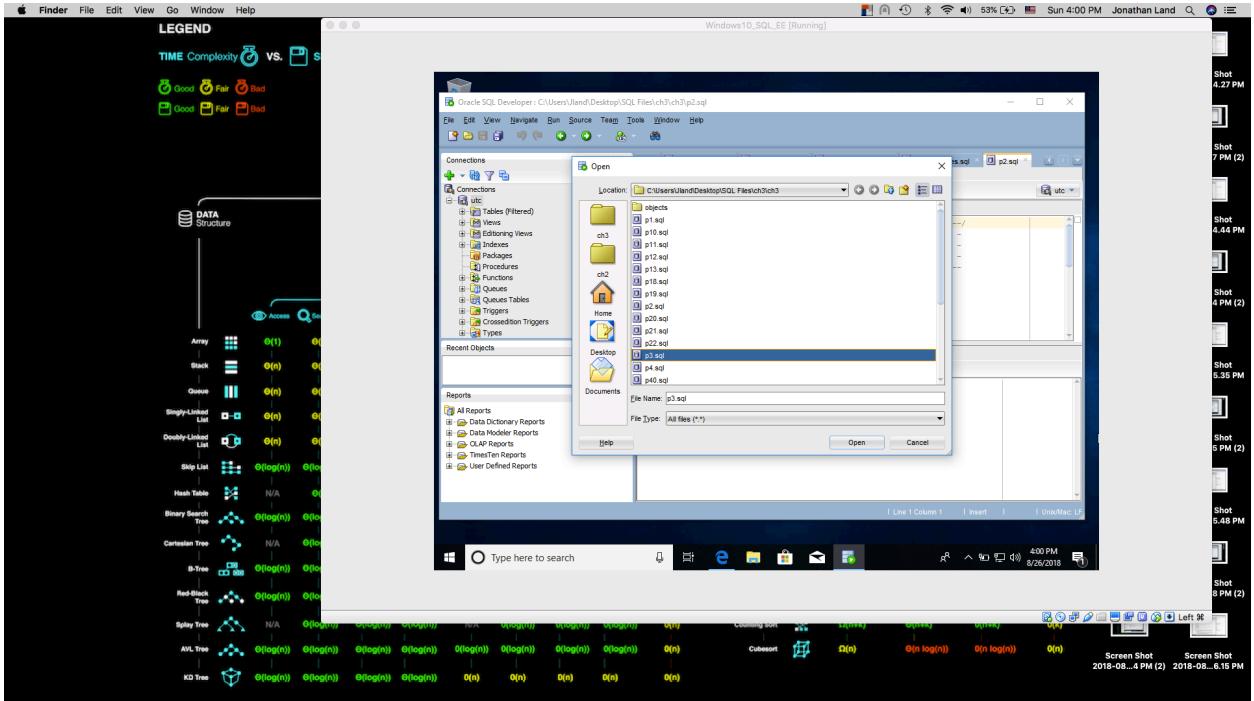
2. Execute scripts in plsql.ch3.zip including *p2.sql*, *p3.sql*, *p4.sql*, *p5.sql*, *p6.sql* by typing: **start p2**

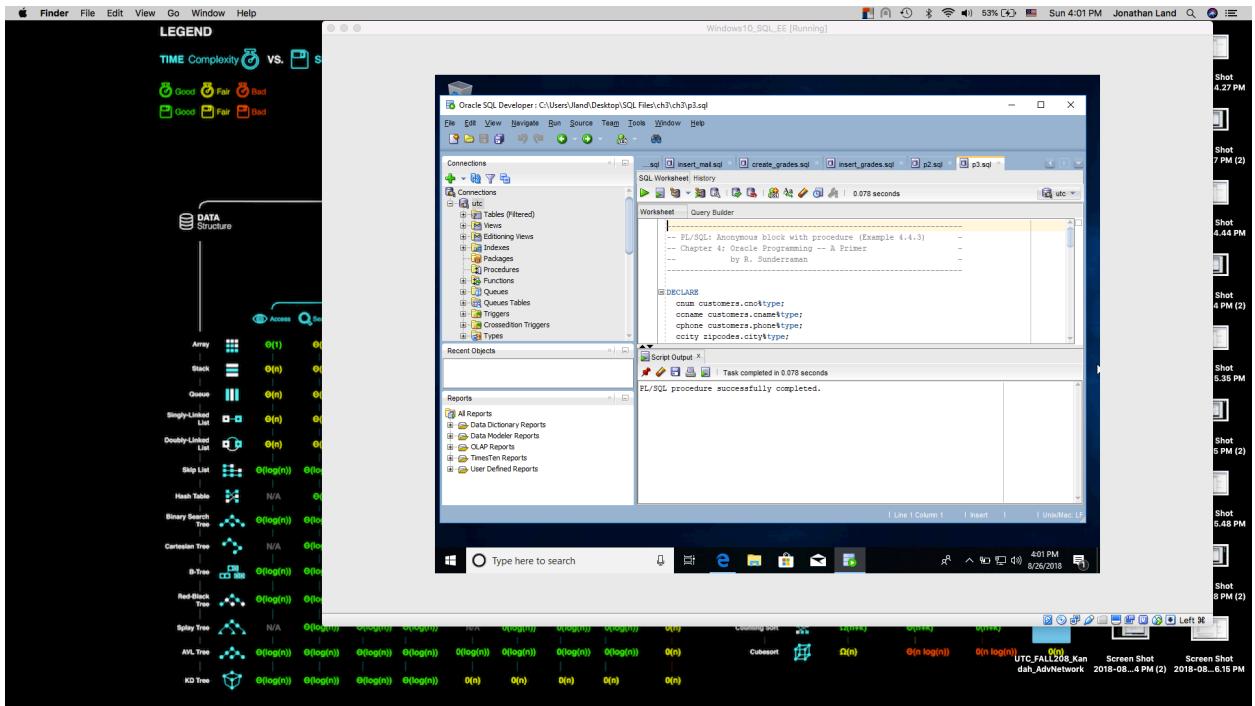
 - a. *p2.sql* gets all the rows from the customers table and prints the names of the customers on the screen.



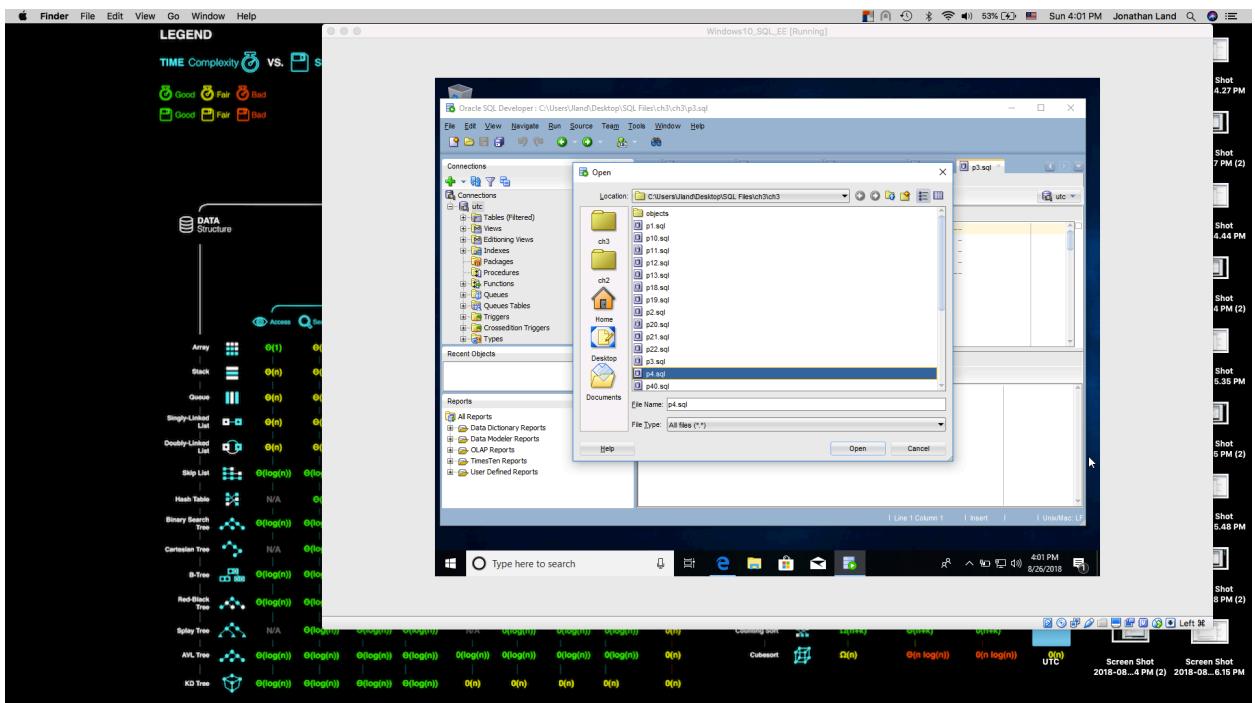


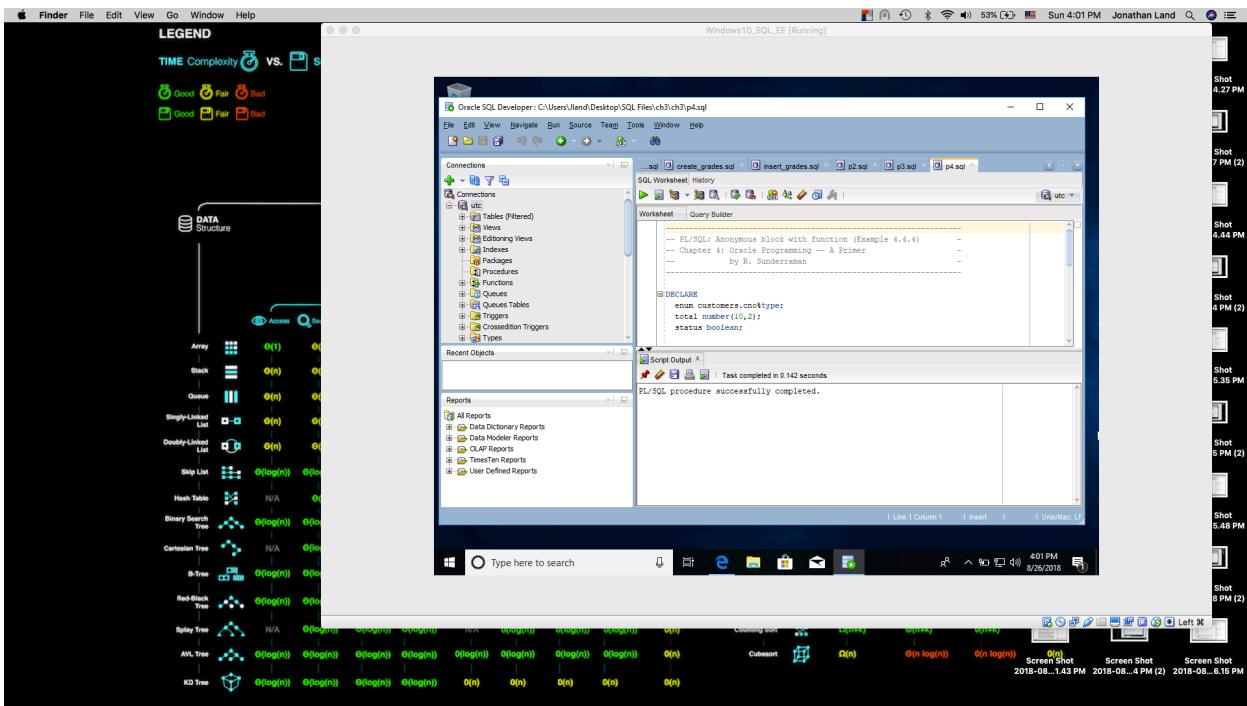
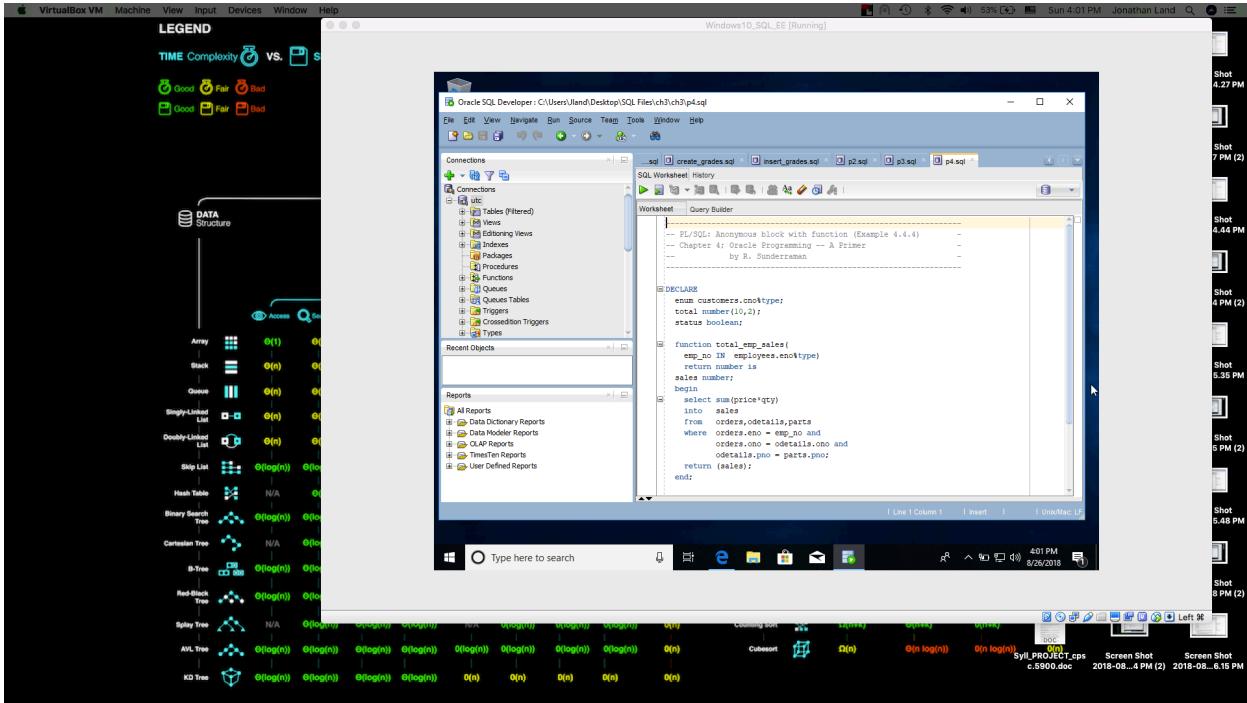
- b. p3.sql contains a procedure that accepts as input a customer number and returns the name, phone, and city values for that customers. If such a customer is not found, status is returned as false; otherwise, status is returned as true.



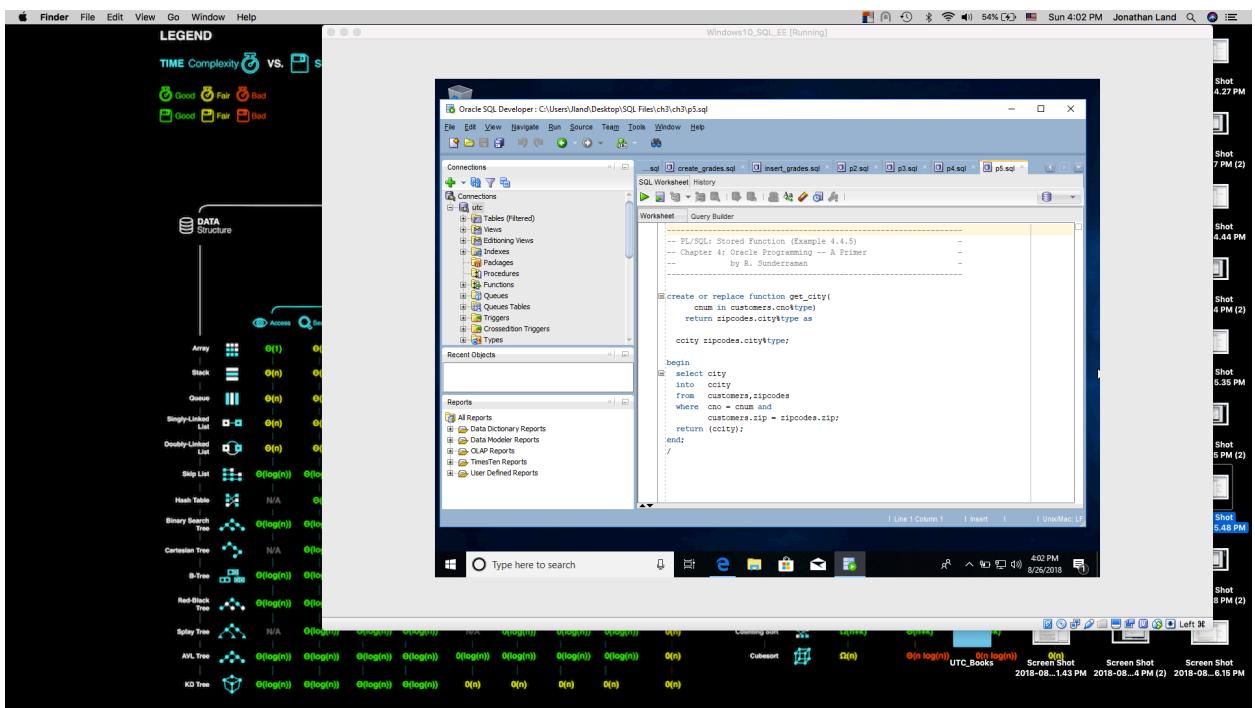
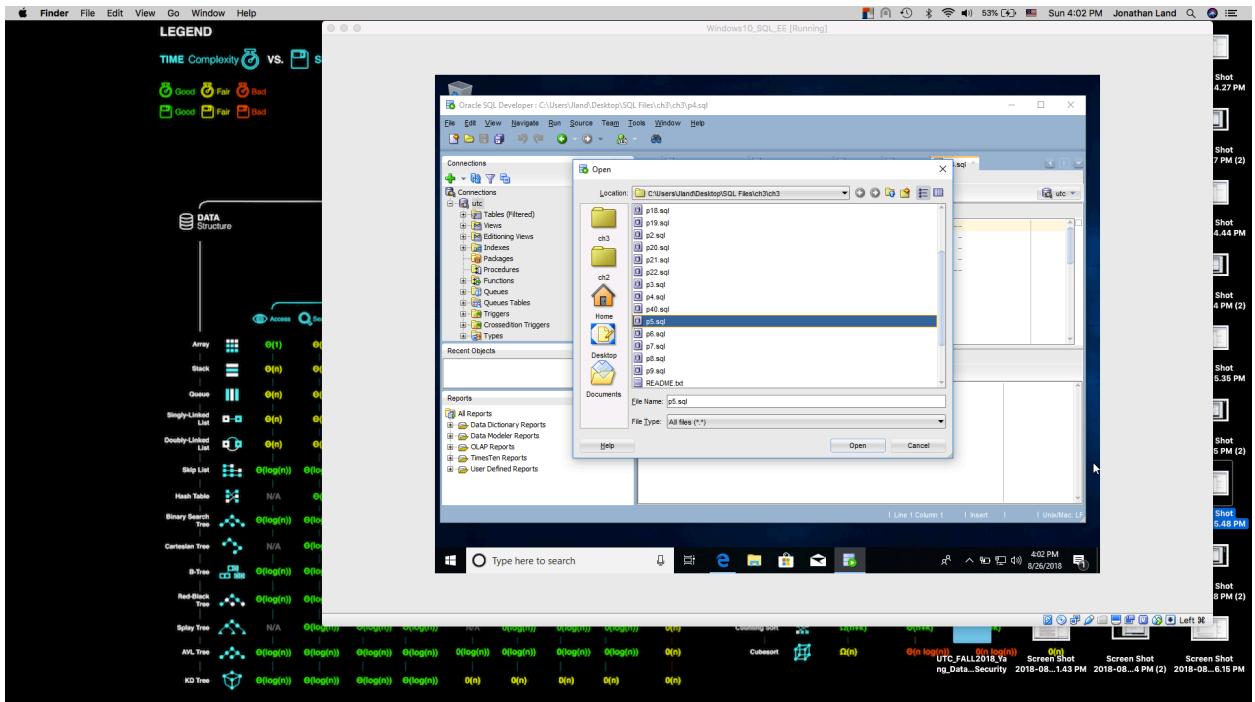


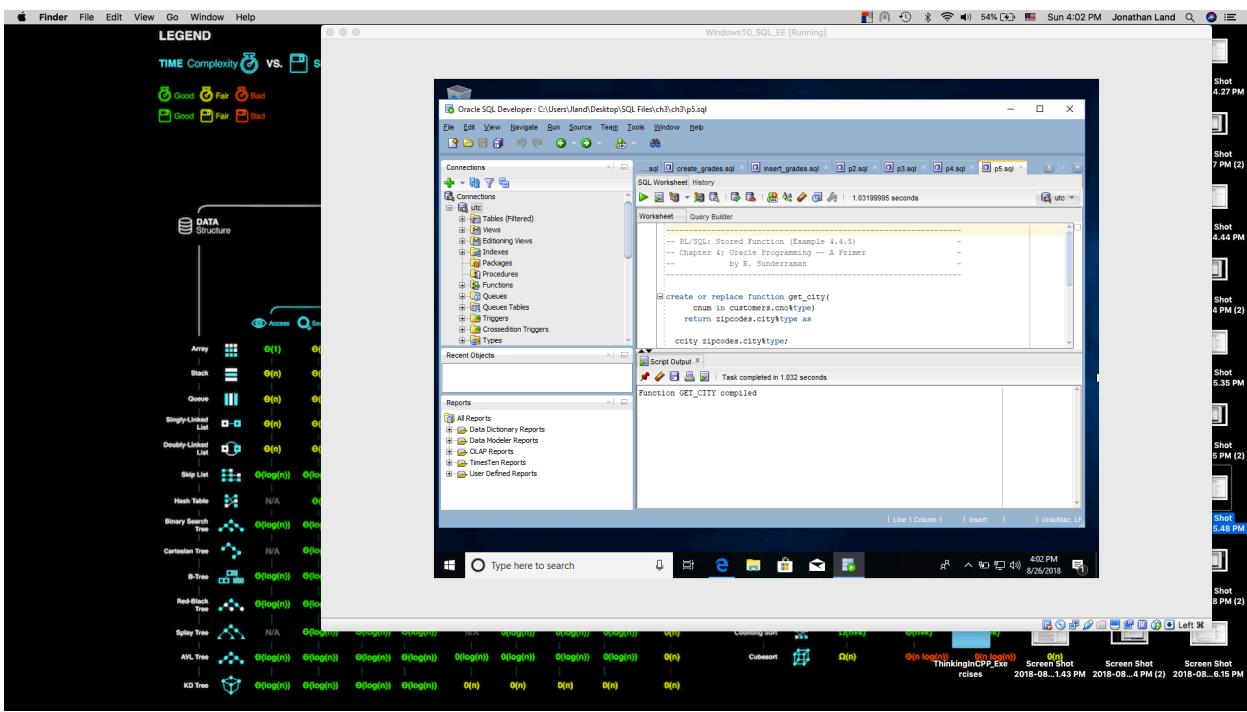
- c. p4.sql contains a function declaration. Given an employee number, the function computes and returns the total sales for that employee.



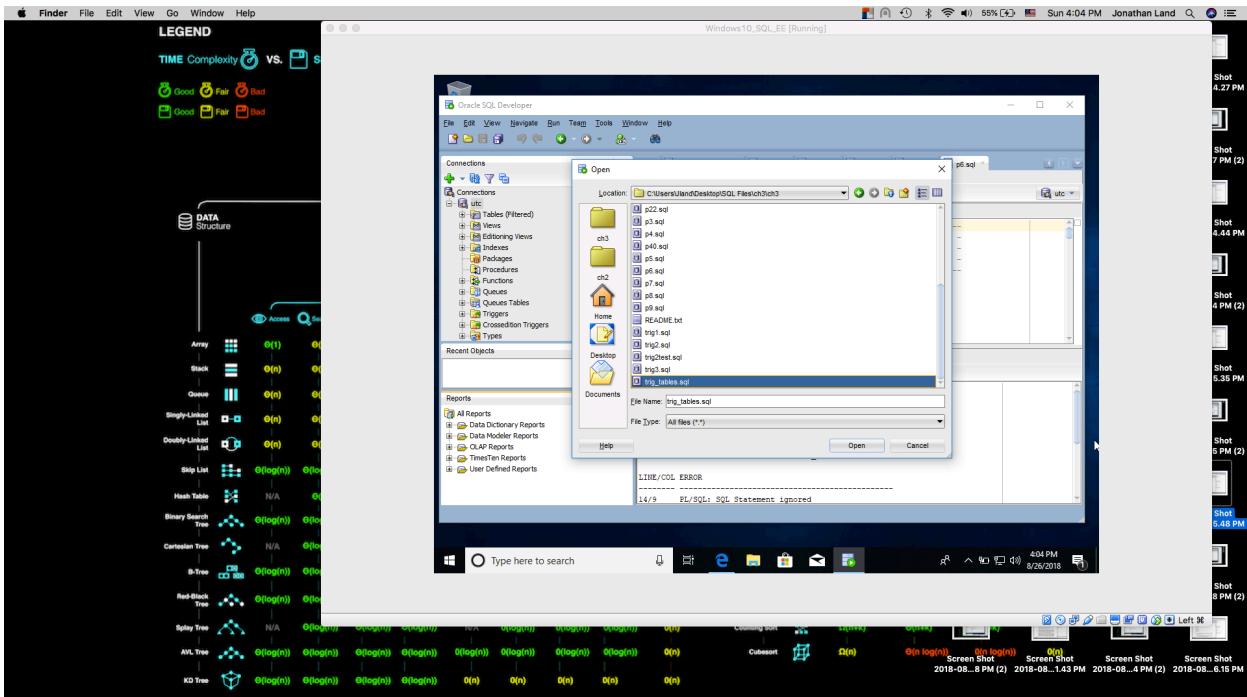


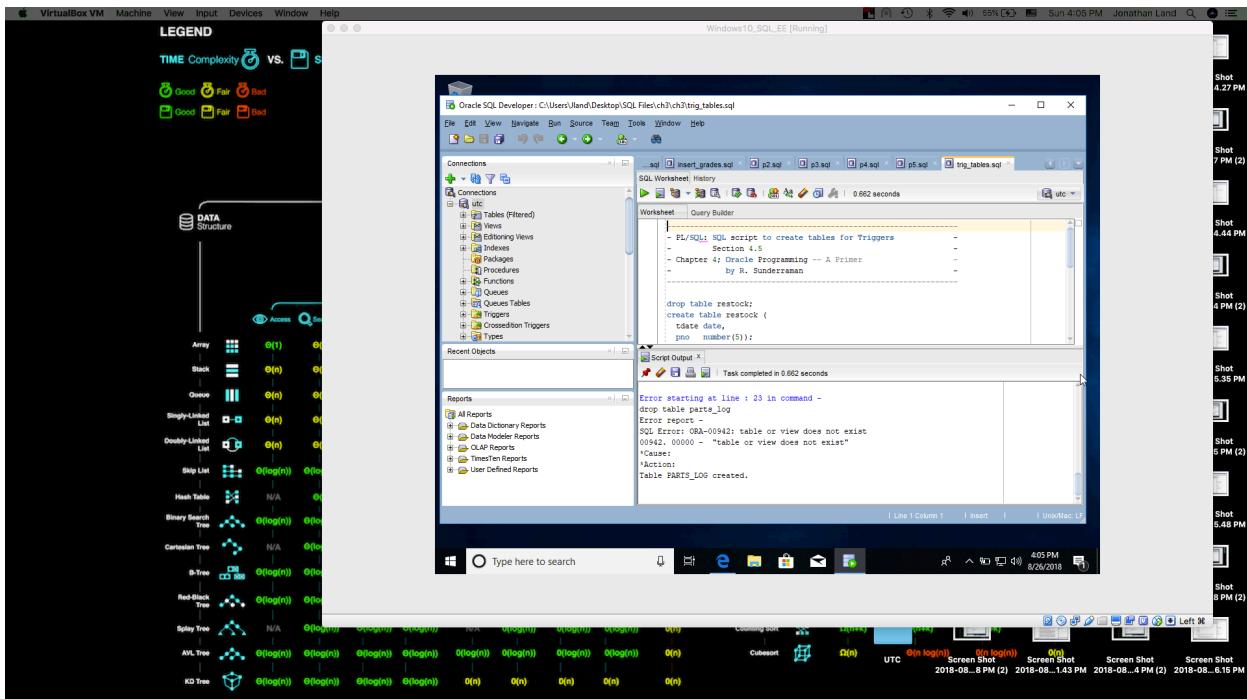
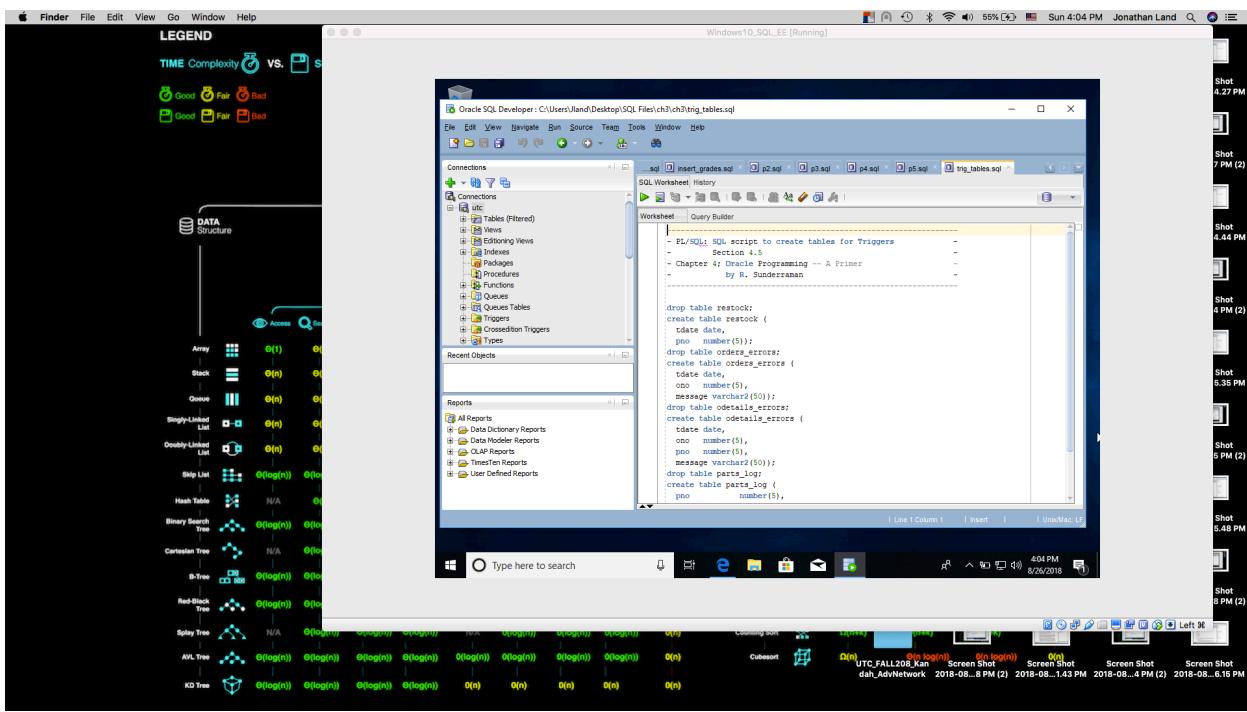
- d. P5.sql contains stored function that takes a customer number as input and returns the city in which the customer lives. A stored procedure can be invoked from various environments, including an SQL statement, another stored function or procedure, an anonymous block or a procedure or function defined within it, an embedded SQL program, a database trigger, and Oracle tools. The stored function get_city can be called by typing: **select cno, cname, get_city(cno);**

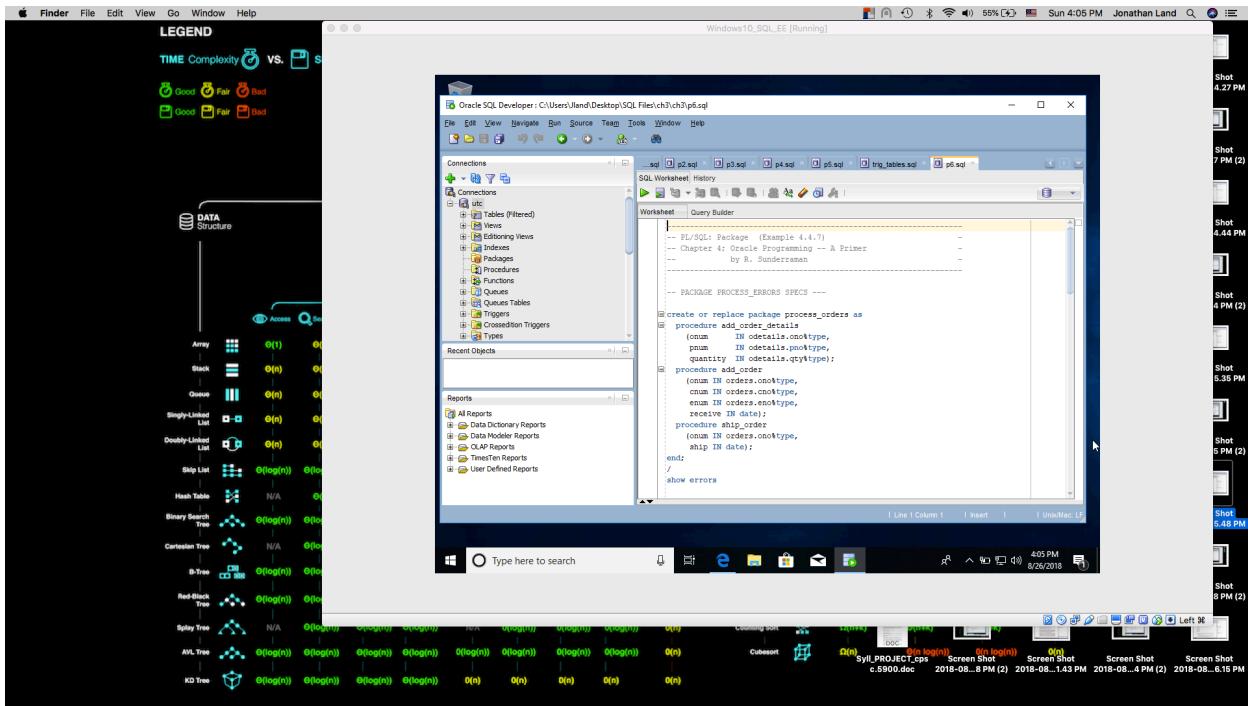




- e. p6.sql defines a package called process_orders containing three procedures. add_order takes as input an order number, customer number, employee number, and received date and tries to insert a new row in the orders table. add_order_details receives as input an order number, part number, and quantity and attempts to add a row corresponding to the input in the odetails table. ship_order takes as input an order number and a shipped date and tries to update the shipped value for the order. Note: You need run trig_tables in ch03 first before you run p6.sql.

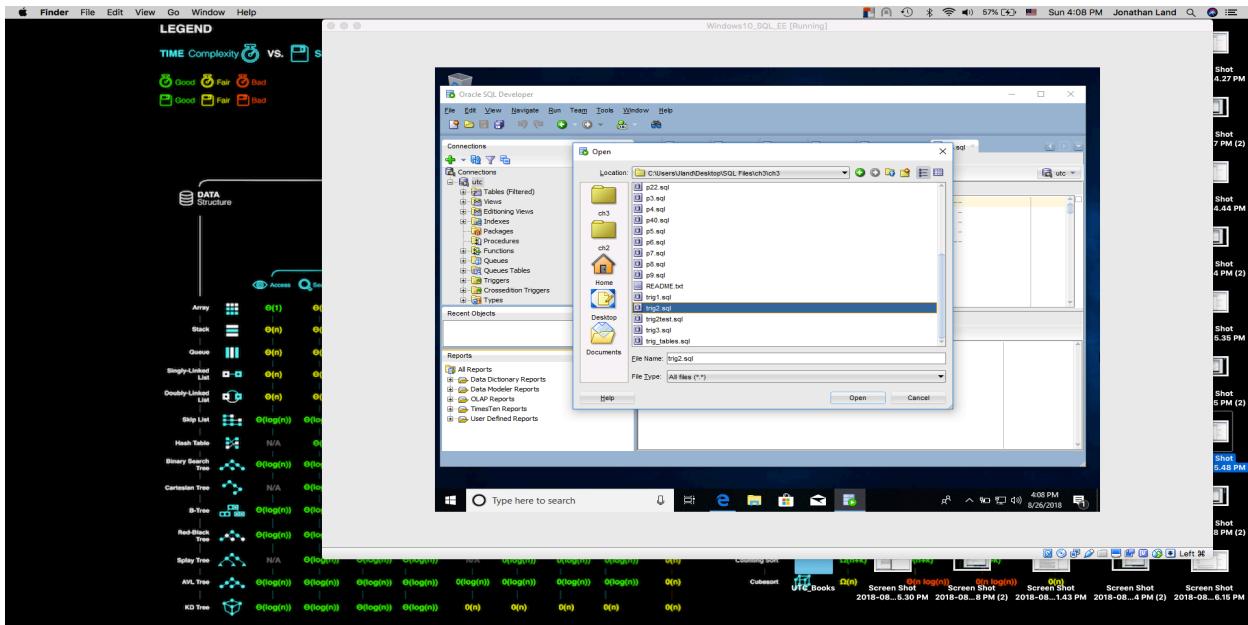


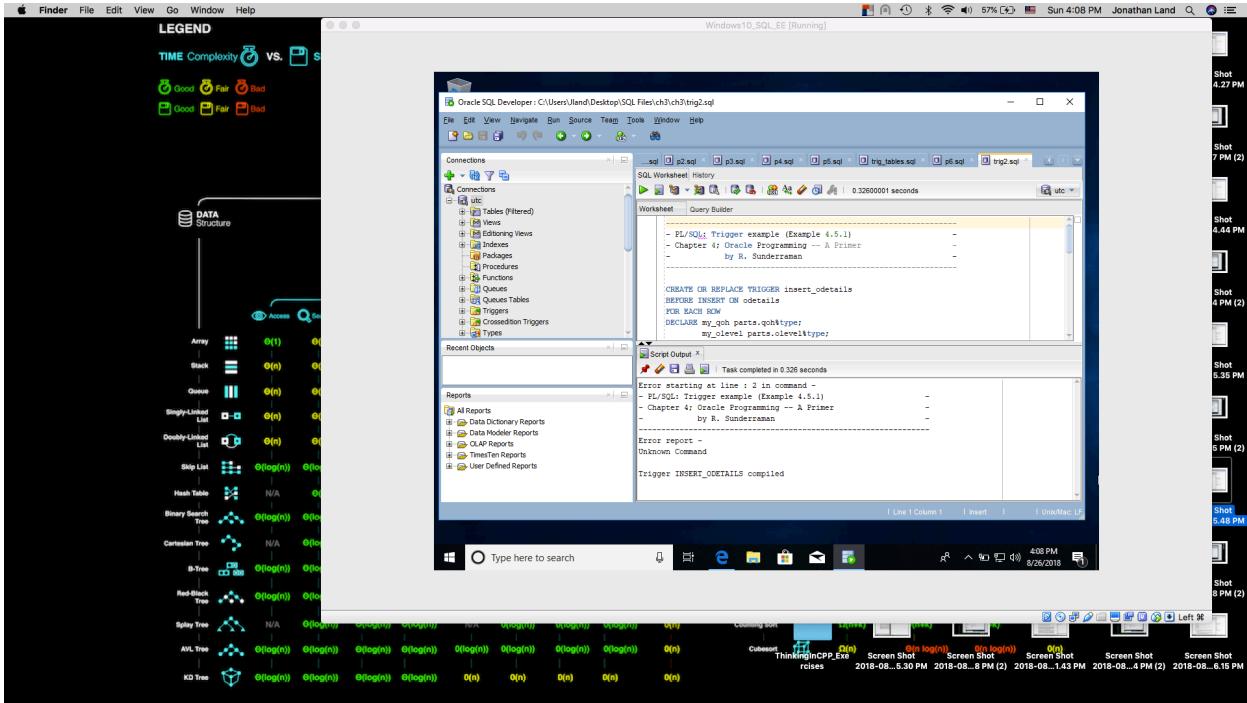




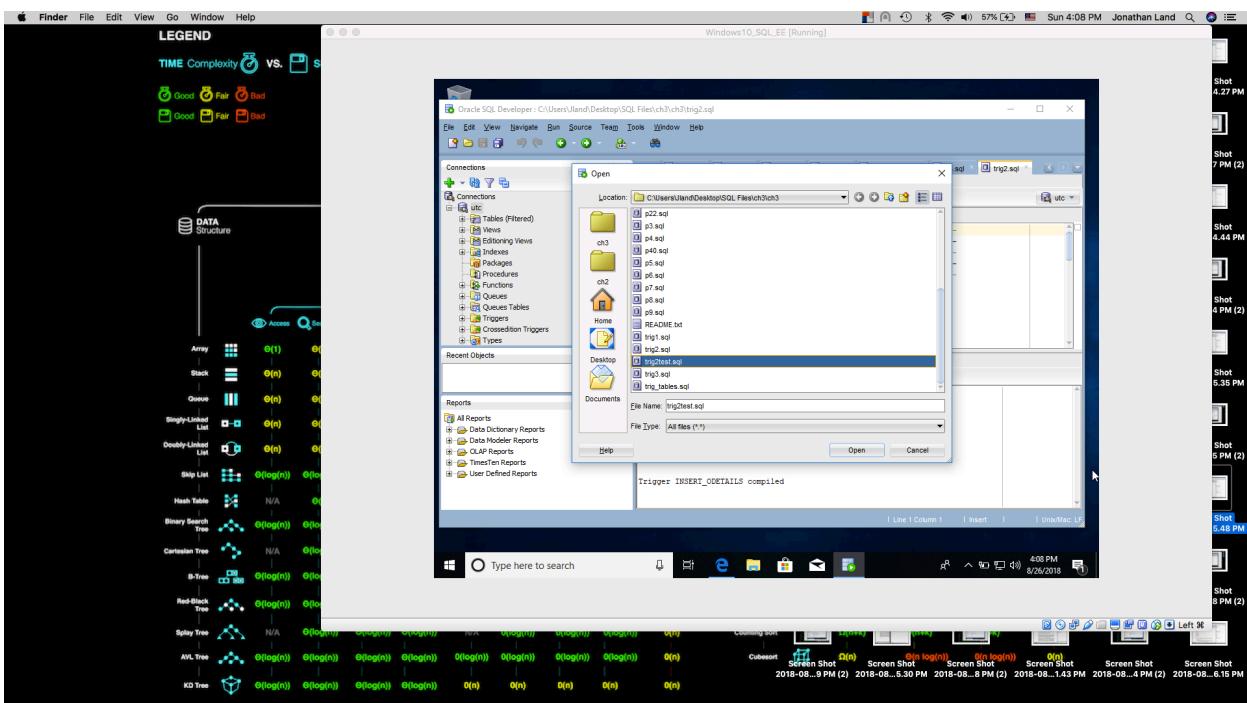
3. Execute `trig2.sql`, `trig2test.sql`, `trig3.sql` in `plsql.ch3.zip`

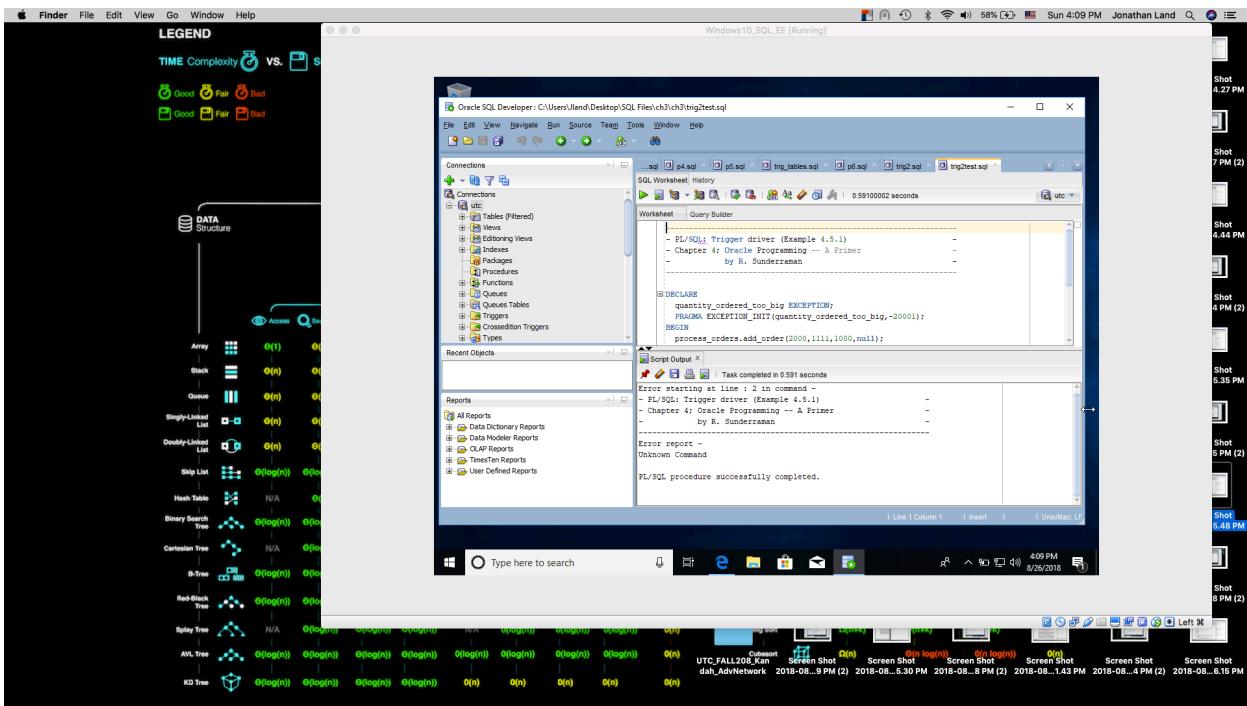
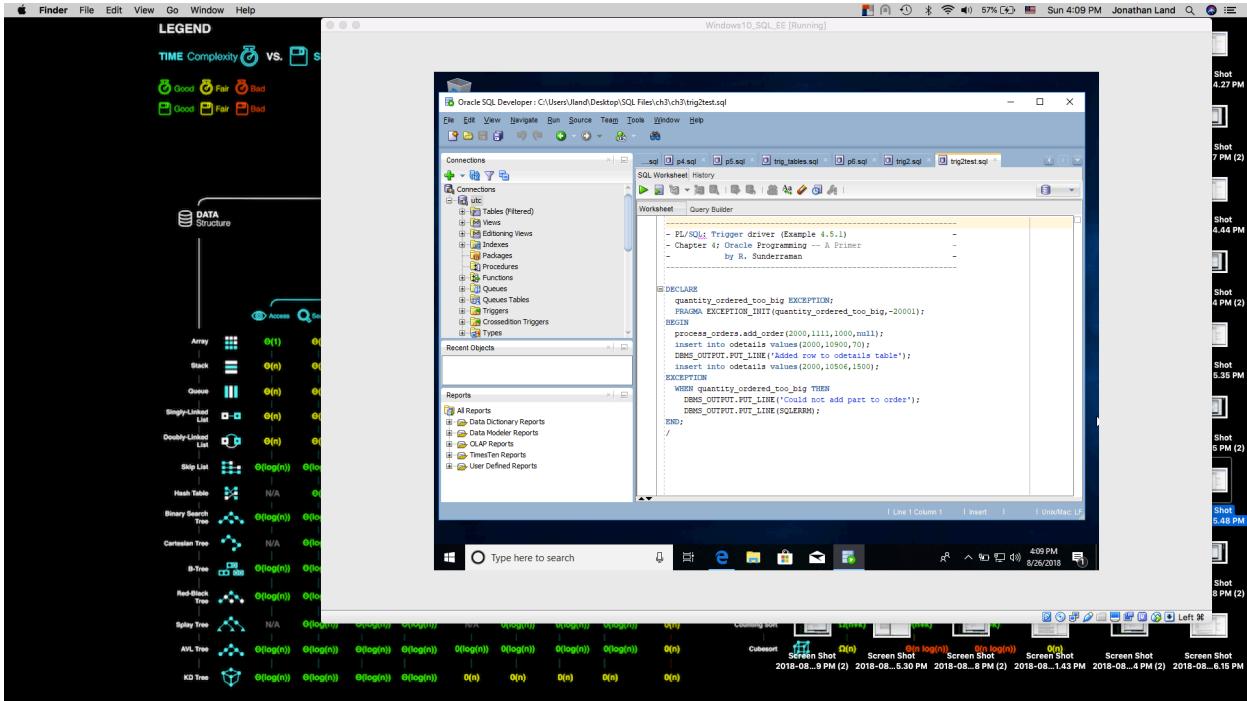
- In `trig2.sql`, a trigger is executed when a row is inserted into the `odetails` table. The trigger checks to see if the quantity ordered is more than the quantity on hand. If it is, an error message is generated, and the row is not inserted. Otherwise, the trigger updates the quantity on hand for the part and checks to see if it has fallen below the reorder level. If it has, it sends a row to the restock table indicating that the part needs to be reordered. You may need to create a `restock` table here, with the first column is `TDATE`, and the second column is `PNO`.





b. trig2test.sql is executed to create a new order.

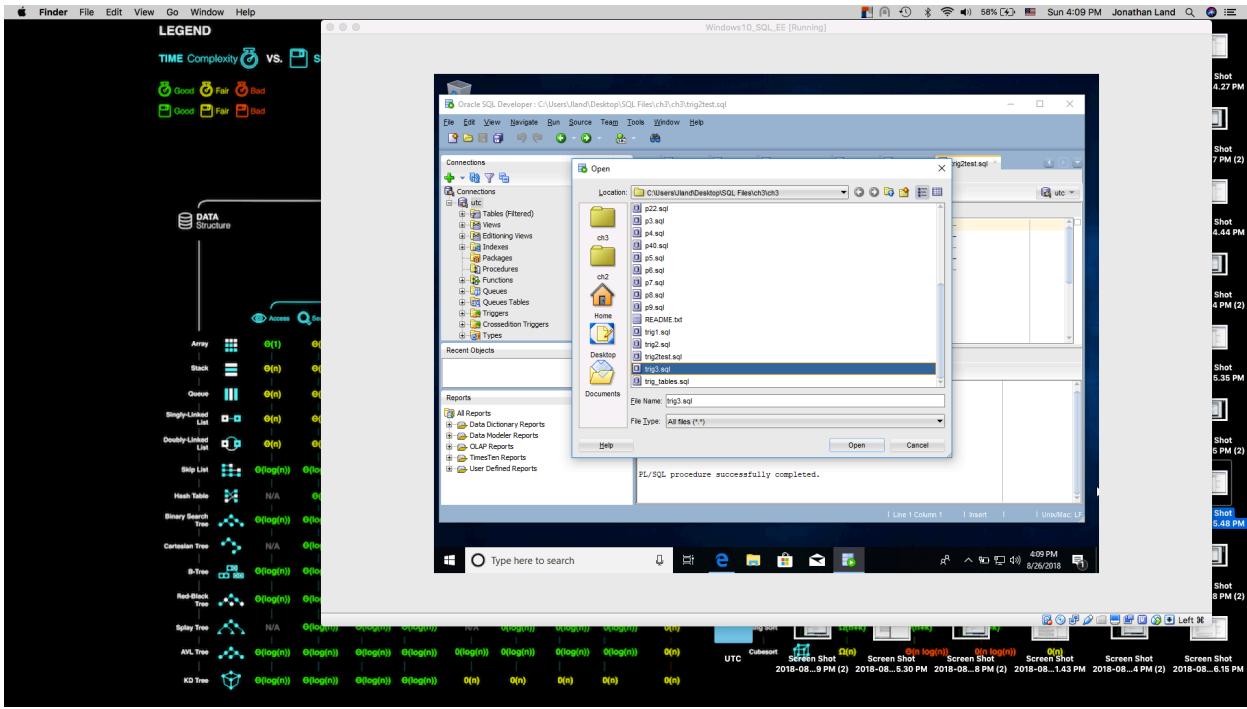




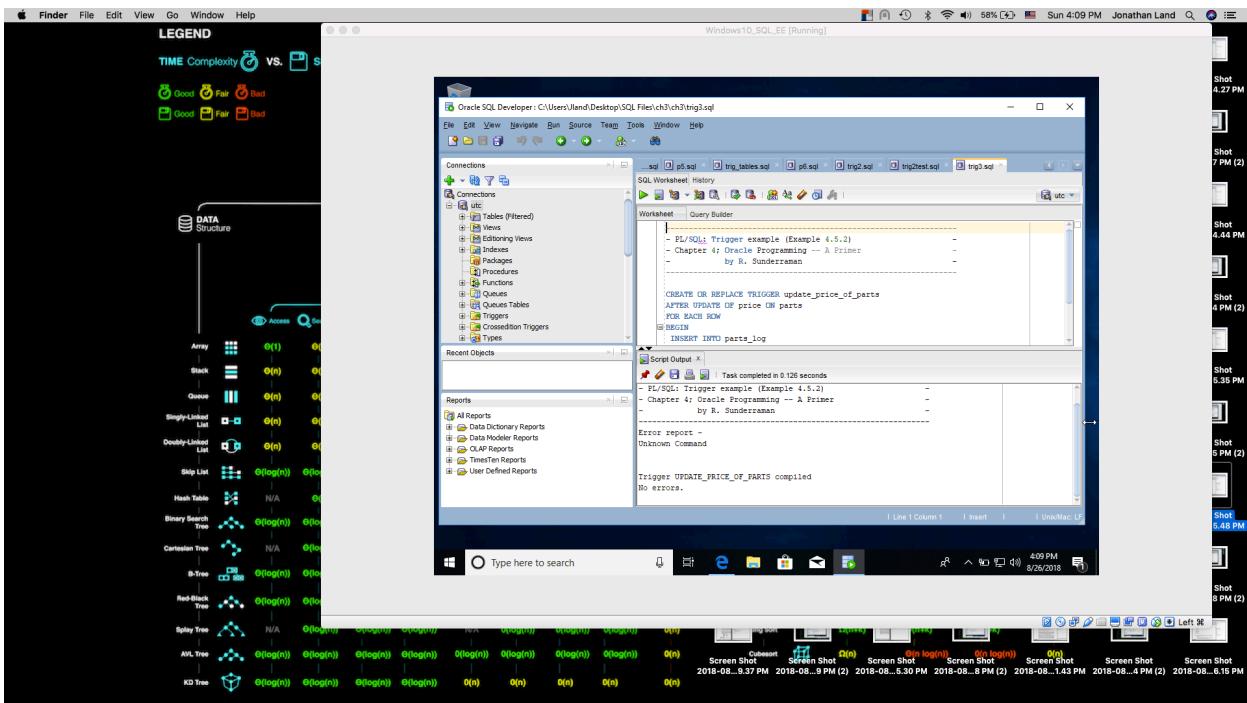
- c. trig3 is defined on the parts table and is triggered when the price column is updated. Each time someone updates the price of a particular part, the trigger makes an entry in a log file of this update along with the userid of the person performing the update and the date the update. The log file is created using
`create table parts_log(`

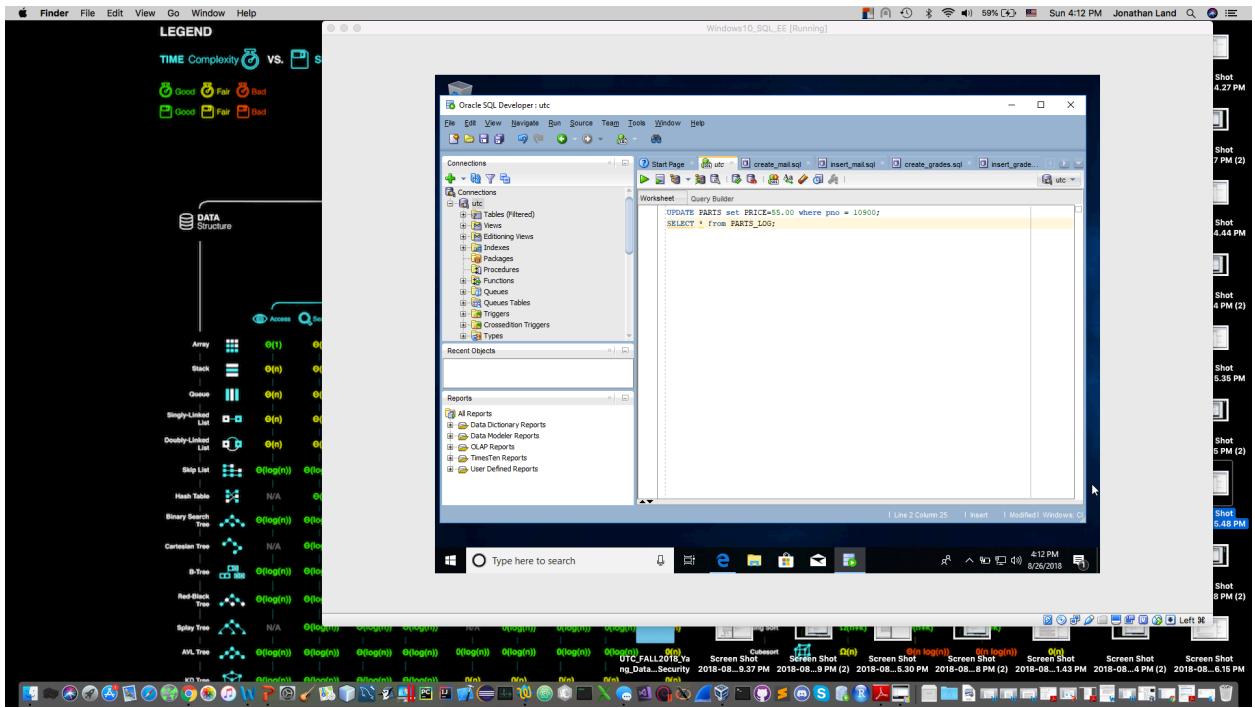
<code>pno</code>	<code>number(5),</code>
<code>username</code>	<code>char(8),</code>

```
update_date date,
old_price number(6,2),
new_price number(6,2));
```

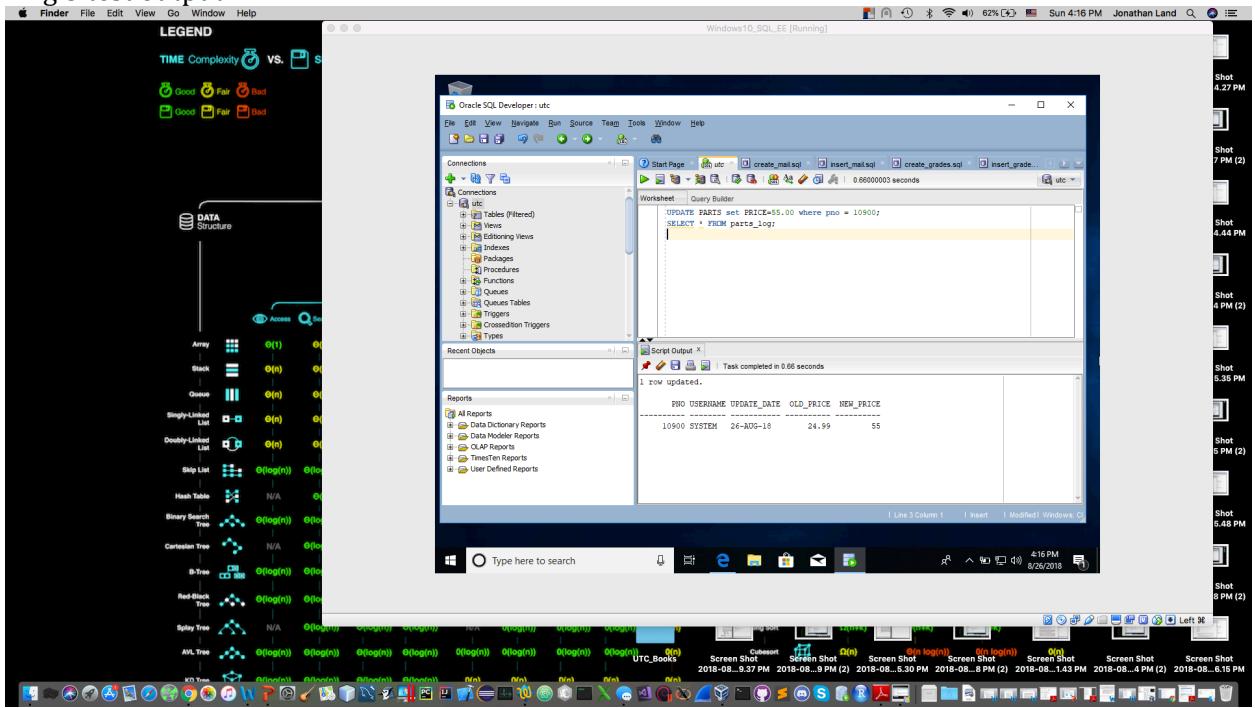


- Test trig3.sql by typing:
update parts set price=55.00 where pno=10900;
select * from parts_log;





Trig 3 test output



5. Create Grade book database using scripts of create_grade.sql and insert_grades.sql

Problem one (40 points): Write a PL/SQL procedure that finds the student with the highest overall average in every course. In case there is more than one student tied for the highest average, the procedure should return all of the students. The results should be returned in a PL/SQL table of records, where each record has the field term, line number, course title, student ID, student name, and overall average. Also, write an anonymous

PL/SQL block that makes a call to the procedure and prints the results to the standard output. Refer to p4.sql for some ideas.

For this problem, two sql scripts were created (p2_problem1.sql and p2_blocks.sql [code below]). In p2_problem1.sql the basic structure/outline for getting the average was as follows: (1) create a view (virtual table) called studentRecords so that the data would always be the current; (2) use the sum function to get the average of student data and then use an alias to reference the data; (3) use inner joins and group the student data from different tables, fields, etc.; (4) create a procedure called **maxAverage** and open a query using sys_refcursor output parameter to order and find the maxAverage of the total combined averages of the students (the cursor starts before the first row in the result set and then “begin” is where the executable table starts); (5) select student sid, fname, lname, studMax.linene, studMax.term, and MAXTotal, ctitle **from studentRecords** created earlier and group that data; (6) select everything from studentRecords and order the max averages via student id, student marks, and max total. The course number and title also show in the rows.

As can be seen below, both studentRecords and maxAverage were created below and the top averages are listed.

The screenshot shows the Oracle SQL Developer interface on a Mac OS X desktop. The main window displays a SQL Worksheet with the following code:

```

1 create or replace view studentRecords AS
2 SELECT SUM(score) / COUNT(score) AS Total,enrolls.sid, courses.lineno,
3 components.term, catalog.title
4 FROM ((scores INNER JOIN enrolls ON (scores.sid=enrolls.sid) AND
5 (scores.term=enrolls.term) AND (scores.lineno=enrolls.lineno))
6 INNER JOIN components ON (scores.components.companame) AND
7 (components.lineno=components.sid) AND (scores.lineno=components.term) AND
8 courses.lineno=components.lineno)
9 INNER JOIN courses ON courses.term=components.term AND
10 courses.lineno=components.lineno
11 GROUP BY enrolls.sid, courses.lineno,components.term,catalog.title;
12
13 create or replace procedure maxAverage(prc out sys_refcursor) is
14 begin
15 open prc for

```

The 'Script Output' pane shows the message: "View STUDENTRECORDS created." and "Procedure MAXAVERAGE compiled".

A floating window titled "TIME COMPLEXITY" is overlaid on the screen, showing a diagram of time complexity classes: O(1), O(log n), O(n), O(n log n), and O(n^2). It also includes a section on "SPACE Complexity" with a "Worst" case analysis.

Windows10_SQL_EE [Running] Sun 10:50 PM Jonathan Land

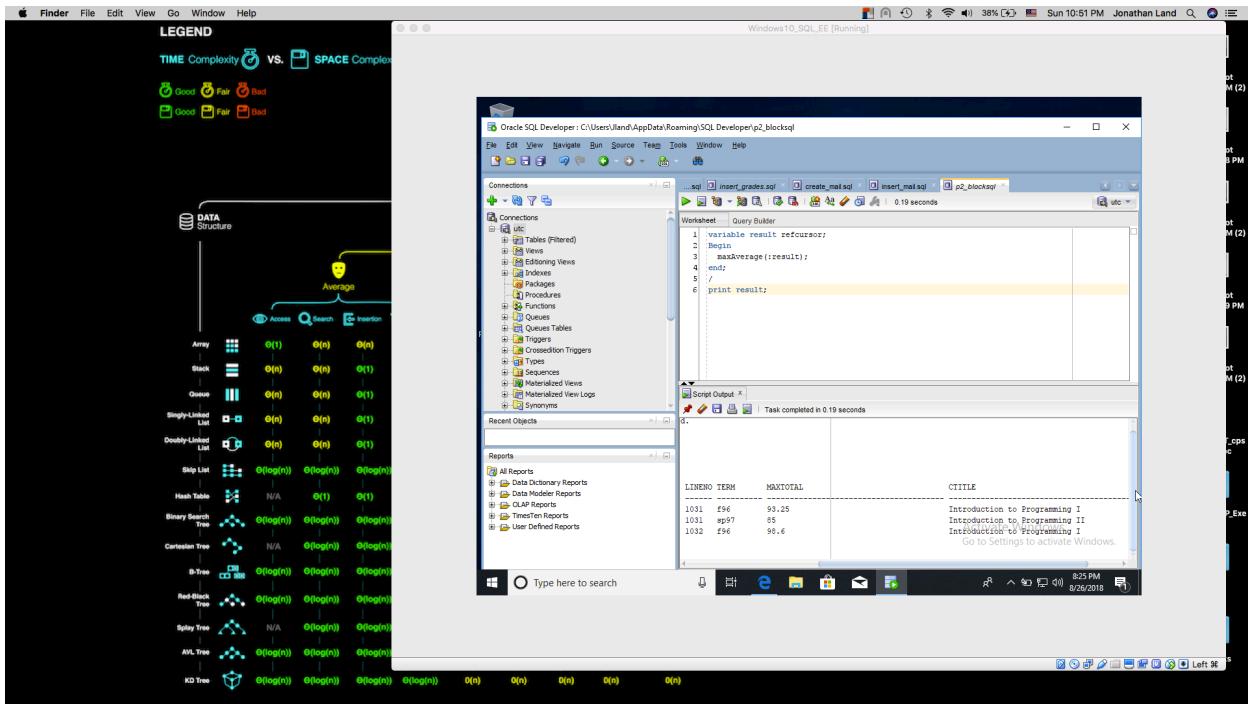
The screenshot shows the Oracle SQL Developer interface with a query window displaying a PL/SQL procedure. The procedure calculates a max average value from a result set. The results table is as follows:

SID	FNAME	LNAME	LINENO	TERM	MAXTOTAL
1111	Nandita	Rajshekhar	1031	f96	93.25
5555	Eld	Yam	1031	sp97	85
1111	Nandita	Rajshekhar	1032	f96	90.6

To the right of the application window, there is a vertical complexity analysis diagram titled "SPACE Complexity". It lists various time complexities from worst to best: $O(n \log n)$, $O(n)$, $O(1), O(\log n)$, $O(1)$, $O(\log(n))$, $O(n)$, $O(1)$, $O(1)$, $O(1)$, $O(1)$, $O(1)$, $O(n)$, $O(n+k)$, $O(k)$, and $O(n)$.

Windows10_SQL_EE [Running] Sun 10:50 PM Jonathan Land

This screenshot is identical to the one above, showing the same Oracle SQL Developer session, query results, and the "SPACE Complexity" analysis diagram.



Problem 1 PL/SQL Code

p2_problem1.sql (maxAverage)

```

create or replace view studentRecords AS
SELECT SUM(points*weight/maxpoints) as Total,enrolls.sid, courses.lineno,
components.term, catalogs.ctitle

FROM (((scores INNER JOIN enrolls ON (scores.sid=enrolls.sid) AND
(scores.term=enrolls.term) AND (scores.lineno=enrolls.lineno))
INNER JOIN components ON (scores.compname=components.compname) AND
(scores.term=components.term) AND (scores.lineno=components.lineno))
INNER JOIN courses ON courses.term=components.term AND
courses.lineno=components.lineno)
INNER JOIN catalogs ON catalogs.cno=courses.cno

GROUP BY enrolls.sid, courses.lineno,components.term,catalogs.ctitle;

create or replace procedure maxAverage(prc out sys_refcursor) is

begin
  open prc for

Select students.sid, fname, lname,studMax.lineno,studMax.term,MAXTotal,ctitle

FROM ((SELECT MAX(Total) as MAXTotal,studentRecords.lineno, studentRecords.term

FROM studentRecords

GROUP BY studentRecords.lineno,studentRecords.term)

```

```

studMax INNER JOIN

(SELECT * FROM studentRecords) studMark ON (MAXTotal=Total) and
(studMax.lineno=studMark.lineno)AND (studMax.term=studMark.term))

INNER JOIN students ON students.sid=studMark.sid;
end;

```

p2_blocks.sql (tester)

```

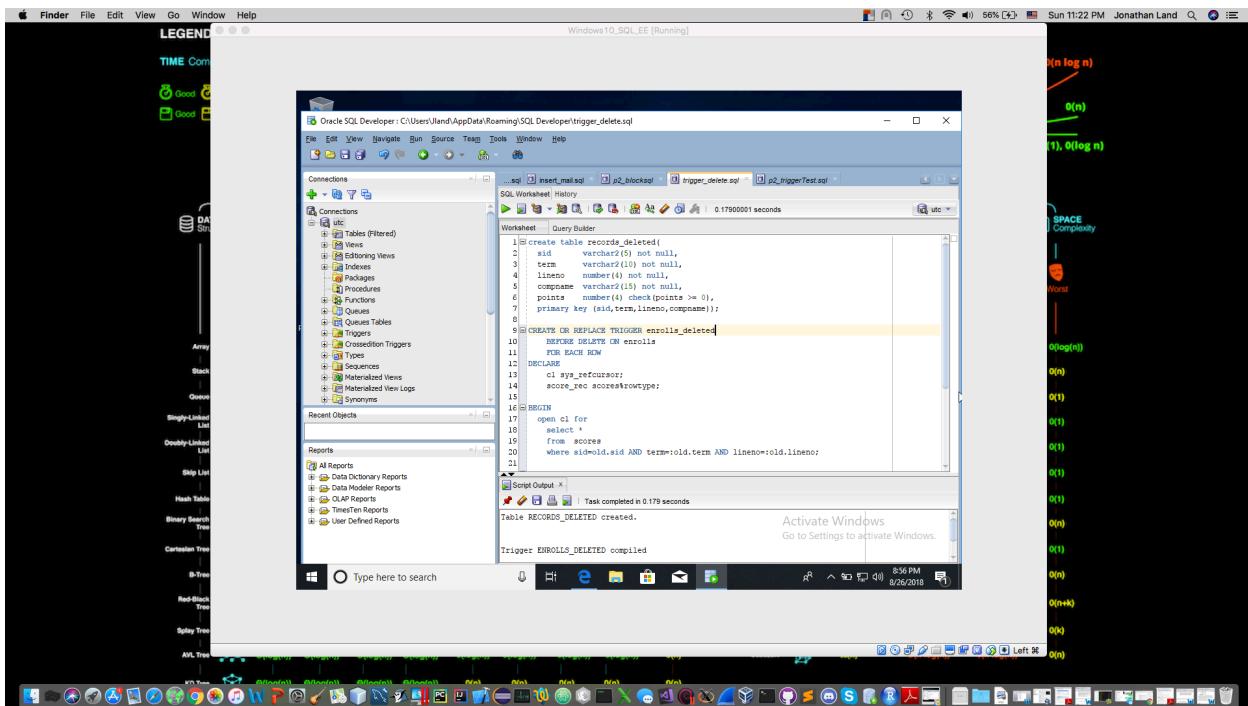
variable result refcursor;
Begin
    maxAverage(:result);
end;

/
print result;

```

Problem two (40 points): Write a trigger that fires when a row is deleted from the enrolls table. The trigger should record the dropped student's scores in a temporary table, called deleted_scores, and cascade the deletes to ensure that the referential integrity constraints are maintained.

For each record we are doing an insert and a delete and will do this until the table is empty (using sys_refcursor). Instead of deleting a row, the trigger kicks off, copies the deleted record in records deleted and then it *actually* deletes it (screenshots and code below)



Windows10_SQL_EE [Running]

```

1: 12  DECLAS
13:   cl sys_refcursor;
14:   score_rec score%rowtype;
15:
16: 13 BEGIN
17:   open cl for
18:   select *
19:   from scores
20:   where sid=:old.sid AND term=:old.term AND lineno=:old.lineno;
21:
22: 14 loop;
23:   fetch cl into score_rec;
24:   exit when cl%NOTFOUND;
25:
26: 15 INSERT INTO records_deleted VALUES (score_rec.sid, score_rec.term,
27:                                         score_rec.lineno, score_rec.comname, score_rec.points);
28: 16 close cl;
29:
30: 17 delete from scores WHERE sid=:old.sid AND term=:old.term AND
31:   lineno=:old.lineno;
32:
33: 18 end;

```

Script Output X | Task completed in 0.183 seconds
Object name = "name is already used by an existing object"
Cause:
Action:
Trigger ENROLLS_DELETED compiled

Windows10_SQL_EE [Running]

```

1: 1 delete from enrolls where sid=4444;
2: 2 SELECT * FROM RECORDS_DELETED;

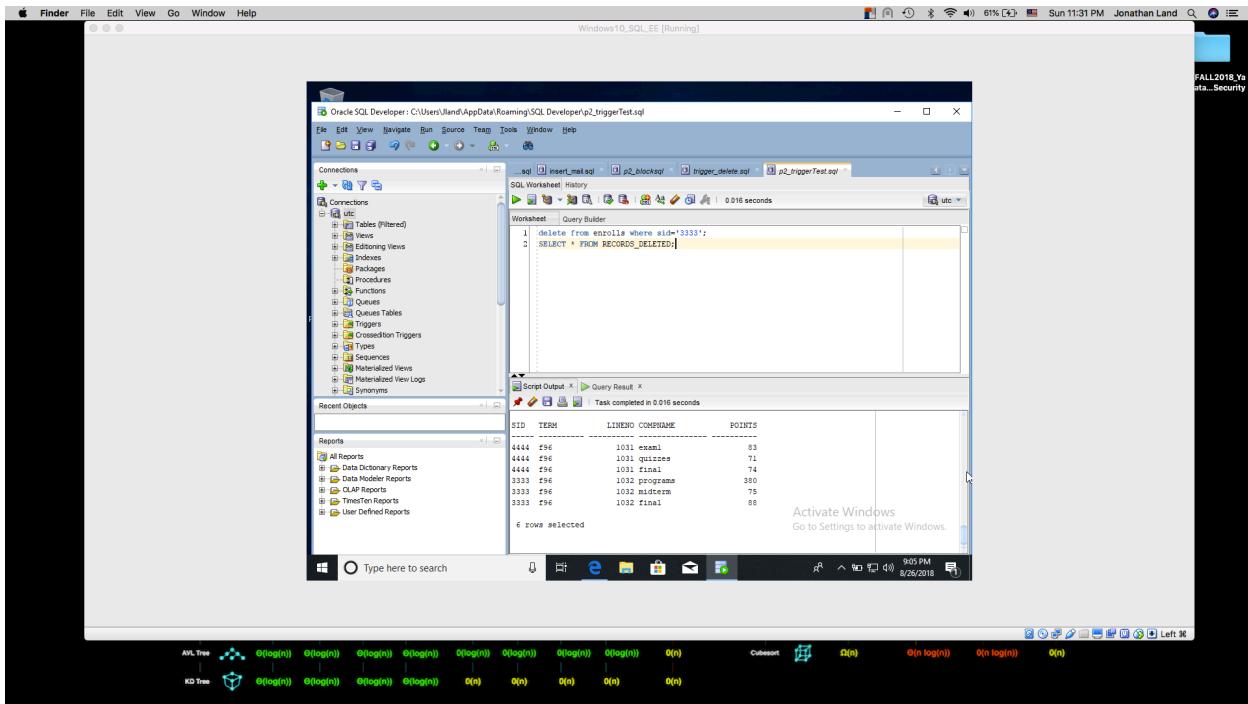
```

Script Output X | Task completed in 0.099 seconds
no rows selected

1 row deleted.

SID	TERM	LINENO	COMPNAME	POINTS
4444	956	1031	exam1	83
4444	956	1031	quizzes	71
4444	956	1031	final	74

Activate Windows
Go to Settings to activate Windows.



Trigger

```

create table records_deleted(
    sid      varchar2(5) not null,
    term     varchar2(10) not null,
    lineno   number(4) not null,
    compname varchar2(15) not null,
    points   number(4) check(points >= 0),
    primary key (sid,term,lineno,compname));

CREATE OR REPLACE TRIGGER enrolls_deleted
    BEFORE DELETE ON enrolls
    FOR EACH ROW
DECLARE
    c1 sys_refcursor;
    score_rec scores%rowtype;

BEGIN
    open c1 for
        select * from scores
        where sid=:old.sid AND term=:old.term AND lineno=:old.lineno;

    loop
        fetch c1 into score_rec;
        exit when c1%NOTFOUND;

        INSERT INTO records_deleted VALUES (score_rec.sid, score_rec.term,
        score_rec.lineno, score_rec.compname, score_rec.points);
    end loop;
END;

```

```
end loop;
close c1;

delete from scores WHERE sid=:old.sid AND term=:old.term AND lineno=:old.lineno;
end;
```

Trigger Tester

```
delete from enrolls where sid='3333';

SELECT * FROM RECORDS_DELETED;
```